

Start coding or generate with AI.

Double-click (or enter) to edit

```
# === BƯỚC 0: Thiết lập & Tải dữ liệu ===
import os, tarfile
import pandas as pd
from sklearn.preprocessing import LabelEncoder

TAR_PATH = "data/hwu.tar.gz"
if os.path.exists(TAR_PATH):
    with tarfile.open(TAR_PATH, "r:gz") as tar:
        tar.extractall("data")
    print("Đã giải nén:", TAR_PATH)
else:
    try:
        from google.colab import files
        print("Upload hwu.tar.gz HOẶC train/val/test (.csv)")
        up = files.upload()
        os.makedirs("data", exist_ok=True)
        for name, content in up.items():
            open(os.path.join("data", name), "wb").write(content)
        for name in up.keys():
            if name.endswith(".tar.gz") or name.endswith(".tgz"):
                with tarfile.open(os.path.join("data", name), "r:gz") as tar:
                    tar.extractall("data")
    except Exception as e:
        print("Không chạy trong Colab hoặc upload thất bại:", e)

data_dir = "data/hwu" if os.path.isdir("data/hwu") else "data"
print("data_dir =", data_dir, "| Files:", os.listdir(data_dir))

train_path = os.path.join(data_dir, "train.csv")
val_path = os.path.join(data_dir, "val.csv")
test_path = os.path.join(data_dir, "test.csv")
assert os.path.exists(train_path) and os.path.exists(val_path) and os.path.exists(test_path), \
    "Không tìm thấy train.csv/val.csv/test.csv trong " + data_dir

df_train = pd.read_csv(train_path)
df_val = pd.read_csv(val_path)
df_test = pd.read_csv(test_path)

def normalize_cols(df):
    cols = {c.lower(): c for c in df.columns}
    text_col = None
    for k in ["text", "utterance", "sentence", "query", "content"]:
        if k in cols: text_col = cols[k]; break
    intent_col = None
    for k in ["intent", "label", "category", "class", "target"]:
        if k in cols: intent_col = cols[k]; break
    if text_col is None: text_col = df.columns[0]
    if intent_col is None: intent_col = df.columns[1]
    return df.rename(columns={text_col:"text", intent_col:"intent"})[["text", "intent"]]

df_train = normalize_cols(df_train)
df_val = normalize_cols(df_val)
df_test = normalize_cols(df_test)

print("Train shape:", df_train.shape)
print("Validation shape:", df_val.shape)
print("Test shape:", df_test.shape)
display(df_train.head())

le = LabelEncoder()
le.fit(pd.concat([df_train["intent"], df_val["intent"], df_test["intent"]], axis=0))

y_train = le.transform(df_train["intent"])
y_val = le.transform(df_val["intent"])
y_test = le.transform(df_test["intent"])
num_classes = len(le.classes_)
print("num_classes:", num_classes, "| ví dụ nhãn:", list(le.classes_)[:10], "...")
```

```
Upload hwu.tar.gz HOẶC train/val/test (.csv)
```

Không có tệp nào được chọn Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving hwu.tar.gz to hwu.tar (6).gz

```
data_dir = data/hwu | Files: ['train_10.csv', 'categories.json', 'val.csv', 'train_5.csv', 'test.csv', 'train.csv']
```

Train shape: (8954, 2)

Validation shape: (1076, 2)

Test shape: (1076, 2)

	text	intent
0	what alarms do i have set right now	alarm_query
1	checkout today alarm of meeting	alarm_query
2	report alarm settings	alarm_query
3	see see for me the alarms that you have set to...	alarm_query
4	is there an alarm for ten am	alarm_query

▼ Nhiệm vụ 1

```
# === NHIỆM VỤ 1: TF-IDF + Logistic Regression ===
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import make_pipeline
from sklearn.metrics import classification_report, f1_score

tfidf_lr_pipeline = make_pipeline(
    TfidfVectorizer(max_features=5000),
    LogisticRegression(max_iter=1000, n_jobs=-1, random_state=42)
)

tfidf_lr_pipeline.fit(df_train["text"], y_train)
y_pred_lr = tfidf_lr_pipeline.predict(df_test["text"])

print("== Classification report: TF-IDF + LR ==")
print(classification_report(y_test, y_pred_lr, target_names=le.classes_, digits=4))
f1_lr = f1_score(y_test, y_pred_lr, average="macro")
print("Macro-F1 (test):", f1_lr)
```

datetime_query	0.7391	0.8947	0.8095	19
email_addcontact	0.7778	0.8750	0.8235	8
email_query	0.8333	0.7895	0.8108	19
email_querycontact	0.9231	0.6316	0.7500	19
email_sendemail	0.8095	0.8947	0.8500	19
general_affirm	1.0000	1.0000	1.0000	19
general_commandstop	1.0000	1.0000	1.0000	19
general_confirm	1.0000	1.0000	1.0000	19
general_dontcare	0.9048	1.0000	0.9500	19
general_explain	1.0000	0.9474	0.9730	19
general_joke	1.0000	1.0000	1.0000	12
general_negate	0.9500	1.0000	0.9744	19
general_praise	0.9500	1.0000	0.9744	19
general_quirky	0.3571	0.2632	0.3030	19
general_repeat	0.9048	1.0000	0.9500	19
iot_cleaning	1.0000	1.0000	1.0000	16
iot_coffee	1.0000	0.9474	0.9730	19
iot_hue_lightchange	0.7500	0.7895	0.7692	19
iot_hue_lighthdim	0.9091	0.8333	0.8696	12
iot_hue_lightoff	0.8947	0.8947	0.8947	19
iot_hue_lighton	0.6667	0.6667	0.6667	3
iot_hue_lightup	1.0000	0.8571	0.9231	14
iot_wemo_off	0.8000	0.8889	0.8421	9

qa_definition	0.8182	0.9474	0.8780	19
qa_factoid	0.4783	0.5789	0.5238	19
qa_maths	0.9231	0.8571	0.8889	14
qa_stock	1.0000	0.9474	0.9730	19
recommendation_events	0.8333	0.7895	0.8108	19
recommendation_locations	0.8095	0.8947	0.8500	19
recommendation_movies	1.0000	1.0000	1.0000	10
social_post	0.9500	1.0000	0.9744	19
social_query	0.8000	0.8889	0.8421	18
takeaway_order	0.8333	0.7895	0.8108	19
takeaway_query	0.8947	0.8947	0.8947	19
transport_query	0.6818	0.7895	0.7317	19
transport_taxi	1.0000	1.0000	1.0000	18
transport_ticket	0.9375	0.7895	0.8571	19
transport_traffic	1.0000	0.9474	0.9730	19
weather_query	0.6190	0.6842	0.6500	19
accuracy		0.8355		1076
macro avg	0.8452	0.8343	0.8353	1076
weighted avg	0.8422	0.8355	0.8351	1076

▼ Nhiệm vụ 2

```
# === NHIỆM VỤ 2: Word2Vec Avg + Dense ===
import numpy as np
from gensim.models import Word2Vec
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping

sentences = [str(s).split() for s in df_train["text"]]
w2v_dim = 100
w2v = Word2Vec(sentences=sentences, vector_size=w2v_dim, window=5, min_count=1, workers=4, seed=42)

def sentence_to_avg_vector(text, model, dim=100):
    toks = str(text).split()
    vecs = [model.wv[t] for t in toks if t in model.wv]
    if not vecs:
        return np.zeros(dim, dtype="float32")
    return np.mean(vecs, axis=0).astype("float32")

def to_matrix(texts, model, dim=100):
    return np.vstack([sentence_to_avg_vector(t, model, dim) for t in texts])

X_train_avg = to_matrix(df_train["text"], w2v, w2v_dim)
X_val_avg = to_matrix(df_val["text"], w2v, w2v_dim)
X_test_avg = to_matrix(df_test["text"], w2v, w2v_dim)

model_avg = Sequential([
    Dense(128, activation='relu', input_shape=(w2v_dim,)),
    Dropout(0.5),
    Dense(num_classes, activation='softmax')
])
model_avg.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

es = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True, verbose=0)
_ = model_avg.fit(X_train_avg, y_train, validation_data=(X_val_avg, y_val),
                   epochs=50, batch_size=64, callbacks=[es], verbose=0)

y_pred_avg = np.argmax(model_avg.predict(X_test_avg, verbose=0), axis=1)
print("== Classification report: W2V-Avg + Dense ==")
print(classification_report(y_test, y_pred_avg, target_names=le.classes_, digits=4))
from sklearn.metrics import f1_score
f1_avg = f1_score(y_test, y_pred_avg, average="macro")
print("Macro-F1 (test):", f1_avg)
```

	0.3555	0.3198	0.3243	17
general_repeat	0.3555	0.3198	0.3243	17
iot_cleaning	0.4375	0.4375	0.4375	16
iot_coffee	0.3810	0.4211	0.4000	19
iot_hue_lightchange	0.4091	0.4737	0.4390	19
iot_hue_lightdim	0.4000	0.1667	0.2353	12
iot_hue_lightoff	0.3208	0.8947	0.4722	19
iot_hue_lighton	0.0000	0.0000	0.0000	3
iot_hue_lightup	0.0000	0.0000	0.0000	14
iot_wemo_off	0.0000	0.0000	0.0000	9
iot_wemo_on	0.3333	0.2857	0.3077	7
lists_createoradd	0.2143	0.4737	0.2951	19
lists_query	0.0000	0.0000	0.0000	19
lists_remove	0.2963	0.4211	0.3478	19
music_likeness	0.0000	0.0000	0.0000	18
music_query	0.1667	0.1053	0.1290	19
music_settings	0.0000	0.0000	0.0000	7
news_query	0.0870	0.1053	0.0952	19
play_audiobook	0.0833	0.0526	0.0645	19
play_game	0.1538	0.1053	0.1250	19
play_music	0.0000	0.0000	0.0000	19
play_podcasts	0.0000	0.0000	0.0000	19
play_radio	0.0769	0.0526	0.0625	19
qa_currency	0.2000	0.0526	0.0833	19
qa_definition	0.0000	0.0000	0.0000	19
qa_factoid	0.0857	0.1579	0.1111	19
qa_maths	0.0000	0.0000	0.0000	14
qa_stock	0.0000	0.0000	0.0000	19
recommendation_events	0.1607	0.4737	0.2400	19
recommendation_locations	0.0000	0.0000	0.0000	19
recommendation_movies	0.0000	0.0000	0.0000	10
social_post	0.0000	0.0000	0.0000	19
social_query	0.0000	0.0000	0.0000	18
takeaway_order	0.1667	0.1053	0.1290	19
takeaway_query	0.5000	0.0526	0.0952	19
transport_query	0.0000	0.0000	0.0000	19
transport_taxi	0.0000	0.0000	0.0000	18
transport_ticket	0.1690	0.6316	0.2667	19
transport_traffic	0.0000	0.0000	0.0000	19
weather_query	0.0000	0.0000	0.0000	19
accuracy		0.2035	1076	
macro avg	0.1667	0.1891	0.1490	1076
weighted avg	0.1646	0.2035	0.1557	1076

Macro-F1 (test): 0.1490183362996796
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for labels with zero samples
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for labels with zero samples
_WARN_PRF(average, modifier, f"{metric.capitalize()} is", len(result))

▼ Nhiệm vụ 3

```
# === NHIỆM VỤ 3: Embedding (pre-trained) + LSTM ===
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.layers import Embedding, LSTM
from tensorflow.keras.models import Sequential

max_words = 20000
oov_tok = "<UNK>"
tokenizer = Tokenizer(num_words=max_words, oov_token=oov_tok)
tokenizer.fit_on_texts(df_train["text"])

def to_pad(texts, tok, max_len=50):
    seqs = tok.texts_to_sequences(texts)
    return pad_sequences(seqs, maxlen=max_len, padding='post', truncating='post')

max_len = 50
X_train_pad = to_pad(df_train["text"], tokenizer, max_len)
X_val_pad = to_pad(df_val["text"], tokenizer, max_len)
X_test_pad = to_pad(df_test["text"], tokenizer, max_len)

vocab_size = min(max_words, len(tokenizer.word_index) + 1)
embedding_dim = w2v_dim

embedding_matrix = np.zeros((vocab_size, embedding_dim), dtype="float32")
for word, idx in tokenizer.word_index.items():
    if idx < vocab_size and word in w2v.wv:
        embedding_matrix[idx] = w2v.wv[word]
```

```

lstm_pre = Sequential([
    Embedding(input_dim=vocab_size, output_dim=embedding_dim,
               weights=[embedding_matrix], input_length=max_len, trainable=False),
    LSTM(128, dropout=0.2, recurrent_dropout=0.2),
    Dense(num_classes, activation='softmax')
])
lstm_pre.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

es = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True, verbose=0)
_ = lstm_pre.fit(X_train_pad, y_train, validation_data=(X_val_pad, y_val),
                  epochs=30, batch_size=64, callbacks=[es], verbose=0)

y_pred_pre = np.argmax(lstm_pre.predict(X_test_pad, verbose=0), axis=1)
print("== Classification report: Emb(pretrained) + LSTM ==")
print(classification_report(y_test, y_pred_pre, target_names=le.classes_, digits=4))
f1_pre = f1_score(y_test, y_pred_pre, average="macro")
print("Macro-F1 (test):", f1_pre)

```

general_commandstop	0.2222	0.3158	0.2609	19
general_confirm	0.0000	0.0000	0.0000	19
general_dontcare	0.0306	0.3684	0.0565	19
general_explain	0.0000	0.0000	0.0000	19
general_joke	0.0000	0.0000	0.0000	12
general_negate	0.0820	0.2632	0.1250	19
general_praise	0.1053	0.6316	0.1805	19
general_quirky	0.0000	0.0000	0.0000	19
general_repeat	0.1600	0.4211	0.2319	19
iot_cleaning	0.2222	0.2500	0.2353	16
iot_coffee	0.0000	0.0000	0.0000	19
iot_hue_lightchange	0.2558	0.5789	0.3548	19
iot_hue_lightdim	0.0000	0.0000	0.0000	12
iot_hue_lightoff	0.4444	0.8421	0.5818	19
iot_hue_lighton	0.0000	0.0000	0.0000	3
iot_hue_lightup	0.0000	0.0000	0.0000	14
iot_wemo_off	0.0000	0.0000	0.0000	9
iot_wemo_on	0.0000	0.0000	0.0000	7
lists_createoradd	0.0129	0.1053	0.0230	19
lists_query	0.0000	0.0000	0.0000	19
lists_remove	0.0000	0.0000	0.0000	19
music_likeness	0.0000	0.0000	0.0000	18
music_query	0.0357	0.0526	0.0426	19
music_settings	0.0000	0.0000	0.0000	7
news_query	0.0000	0.0000	0.0000	19
play_audiobook	0.0000	0.0000	0.0000	19
play_game	0.0000	0.0000	0.0000	19
play_music	0.0000	0.0000	0.0000	19
play_podcasts	0.0000	0.0000	0.0000	19
play_radio	0.0000	0.0000	0.0000	19
qa_currency	0.0000	0.0000	0.0000	19
qa_definition	0.0000	0.0000	0.0000	19
qa_factoid	0.0000	0.0000	0.0000	19
qa_maths	0.0000	0.0000	0.0000	14
qa_stock	0.0000	0.0000	0.0000	19
recommendation_events	0.0000	0.0000	0.0000	19
recommendation_locations	0.0000	0.0000	0.0000	19
recommendation_movies	0.0000	0.0000	0.0000	10
social_post	0.0000	0.0000	0.0000	19
social_query	0.0000	0.0000	0.0000	18
takeaway_order	0.0000	0.0000	0.0000	19
takeaway_query	0.0000	0.0000	0.0000	19
transport_query	0.0000	0.0000	0.0000	19
transport_taxi	0.0000	0.0000	0.0000	18
transport_ticket	0.0000	0.0000	0.0000	19
transport_traffic	0.0000	0.0000	0.0000	19
weather_query	0.0000	0.0000	0.0000	19
accuracy		0.0818		1076
macro avg	0.0266	0.0730	0.0362	1076
weighted avg	0.0294	0.0818	0.0402	1076

Macro-F1 (test): 0.03616698778099441
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for labels with zero samples
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for labels with zero samples
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for labels with zero samples

▼ Nhiệm vụ 4

```
# === NHIỆM VỤ 4: Embedding (scratch) + LSTM ===
lstm_scr = Sequential([
    Embedding(input_dim=vocab_size, output_dim=100, input_length=max_len),
    LSTM(128, dropout=0.2, recurrent_dropout=0.2),
    Dense(num_classes, activation='softmax')
])
lstm_scr.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

es = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True, verbose=0)
_ = lstm_scr.fit(X_train_pad, y_train, validation_data=(X_val_pad, y_val),
                  epochs=30, batch_size=64, callbacks=[es], verbose=0)

y_pred_scr = np.argmax(lstm_scr.predict(X_test_pad), verbose=0, axis=1)
print("== Classification report: Emb(scratch) + LSTM ==")
print(classification_report(y_test, y_pred_scr, target_names=le.classes_, digits=4))
f1_scr = f1_score(y_test, y_pred_scr, average="macro")
print("Macro-F1 (test):", f1_scr)
```

general_commandstop	0.5833	0.3684	0.4516	19
general_confirm	0.2917	0.7368	0.4179	19
general_dontcare	0.4412	0.7895	0.5660	19
general_explain	0.5185	0.7368	0.6087	19
general_joke	0.0000	0.0000	0.0000	12
general_negate	0.0000	0.0000	0.0000	19
general_praise	0.3333	0.4211	0.3721	19
general_quirky	0.0000	0.0000	0.0000	19
general_repeat	0.1250	0.0526	0.0741	19
iot_cleaning	0.0000	0.0000	0.0000	16
iot_coffee	0.2759	0.4211	0.3333	19
iot_hue_lightchange	0.1923	0.5263	0.2817	19
iot_hue_lighthdim	0.0000	0.0000	0.0000	12
iot_hue_lightoff	0.0000	0.0000	0.0000	19
iot_hue_lighton	0.0000	0.0000	0.0000	3
iot_hue_lightup	0.0000	0.0000	0.0000	14
iot_wemo_off	0.0000	0.0000	0.0000	9
iot_wemo_on	0.0000	0.0000	0.0000	7
lists_createoradd	0.1923	0.5263	0.2817	19
lists_query	0.0000	0.0000	0.0000	19
lists_remove	0.2963	0.4211	0.3478	19
music_likeness	0.0000	0.0000	0.0000	18
music_query	0.1250	0.1579	0.1395	19
music_settings	0.0000	0.0000	0.0000	7
news_query	0.2333	0.3684	0.2857	19
play_audiobook	0.2500	0.1053	0.1481	19
play_game	0.2188	0.3684	0.2745	19
play_music	0.1429	0.1053	0.1212	19
play_podcasts	0.2500	0.4737	0.3273	19
play_radio	0.2667	0.4211	0.3265	19
qa_currency	0.1667	0.3158	0.2182	19
qa_definition	0.1519	0.6316	0.2449	19
qa_factoid	0.0000	0.0000	0.0000	19
qa_maths	0.0000	0.0000	0.0000	14
qa_stock	0.2333	0.3684	0.2857	19
recommendation_events	0.0000	0.0000	0.0000	19
recommendation_locations	0.0000	0.0000	0.0000	19
recommendation_movies	0.0000	0.0000	0.0000	10
social_post	0.1538	0.2105	0.1778	19
social_query	0.0000	0.0000	0.0000	18
takeaway_order	0.1111	0.0526	0.0714	19
takeaway_query	0.1714	0.3158	0.2222	19
transport_query	0.2000	0.2105	0.2051	19
transport_taxi	0.3750	0.1667	0.2308	18
transport_ticket	0.2000	0.3158	0.2449	19
transport_traffic	0.0000	0.0000	0.0000	19
weather_query	0.0000	0.0000	0.0000	19
accuracy			0.2314	1076
macro avg	0.1293	0.2049	0.1495	1076
weighted avg	0.1458	0.2314	0.1688	1076

Macro-F1 (test): 0.14952501738646845

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined
    _warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined
    _warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined
```

▼ Nhiệm vụ 5

```
# === NHIỆM VỤ 5: So sánh định lượng + Phân tích định tính ===
import numpy as np
import pandas as pd
from sklearn.metrics import f1_score, classification_report

loss_avg, _ = model_avg.evaluate(X_test_avg, y_test, verbose=0)
loss_pre, _ = lstm_pre.evaluate(X_test_pad, y_test, verbose=0)
loss_scr, _ = lstm_scr.evaluate(X_test_pad, y_test, verbose=0)

summary = pd.DataFrame({
    "Pipeline": [
        "TF-IDF + Logistic Regression",
        "Word2Vec (Avg) + Dense",
        "Embedding (Pre-trained) + LSTM",
        "Embedding (Scratch) + LSTM"
    ],
    "F1-macro (test)": [f1_lr, f1_avg, f1_pre, f1_scr],
    "Test Loss": [None, loss_avg, loss_pre, loss_scr]
}).sort_values("F1-macro (test)", ascending=False).reset_index(drop=True)

display(summary)

hard_texts = [
    "can you remind me to not call my mom",
    "is it going to be sunny or rainy tomorrow",
    "find a flight from new york to london but not through paris"
]

def predict_all(texts):
    pred_lr = tfidf_lr_pipeline.predict(texts)
    Xavg = np.vstack([sentence_to_avg_vector(t, w2v, w2v_dim) for t in texts])
    pred_avg = np.argmax(model_avg.predict(Xavg, verbose=0), axis=1)
    seqs = tokenizer.texts_to_sequences(texts)
    Xpad = pad_sequences(seqs, maxlen=max_len, padding='post', truncating='post')
    pred_pre = np.argmax(lstm_pre.predict(Xpad, verbose=0), axis=1)
    pred_scr = np.argmax(lstm_scr.predict(Xpad, verbose=0), axis=1)
    return pred_lr, pred_avg, pred_pre, pred_scr

pred_lr, pred_avg, pred_pre, pred_scr = predict_all(hard_texts)

true_labels = []
for t in hard_texts:
    match = df_test[df_test["text"].str.lower() == t.lower()]
    if len(match) > 0:
        true_labels.append(le.transform(match["intent"])[0])
    else:
        true_labels.append(None)

rows = []
for i, text in enumerate(hard_texts):
    y_true = true_labels[i]
    row = {
        "text": text,
        "True Intent": le.classes_[y_true] if y_true is not None else "N/A",
        "TF-IDF + LR": le.classes_[pred_lr[i]],
        "W2V-Avg + Dense": le.classes_[pred_avg[i]],
        "LSTM (pre)": le.classes_[pred_pre[i]],
        "LSTM (scratch)": le.classes_[pred_scr[i]],
    }
    if y_true is not None:
        row["✓ LR"] = "✓" if pred_lr[i] == y_true else "X"
        row["✓ W2V"] = "✓" if pred_avg[i] == y_true else "X"
        row["✓ LSTM-pre"] = "✓" if pred_pre[i] == y_true else "X"
        row["✓ LSTM-scr"] = "✓" if pred_scr[i] == y_true else "X"
    rows.append(row)

qual_df = pd.DataFrame(rows)
display(qual_df)

print("\nNhận xét:")
print("- Các câu có phủ định ('not') hoặc mệnh đề phụ ('or', 'but not through ...') thường khiến mô hình TF-IDF và W2V trung bình vì chúng không nắm được thứ tự và phạm vi của phủ định.")
print("- LSTM (đặc biệt bản pre-trained) có xu hướng hiểu ngữ cảnh tốt hơn, vì trạng thái ẩn theo chuỗi giúp mô hình nhận diện đ")
print("- Nếu LSTM (pre-trained) chính xác hơn bản học từ đầu, điều đó chứng tỏ embedding Word2Vec cung cấp ngữ nghĩa ban đầu hữu
```

	Pipeline	F1-macro (test)	Test Loss			
0	TF-IDF + Logistic Regression	0.835298	NaN			
1	Embedding (Scratch) + LSTM	0.149525	2.941058			
2	Word2Vec (Avg) + Dense	0.149018	3.032118			
3	Embedding (Pre-trained) + LSTM	0.036167	3.603505			
	text	True Intent	TF-IDF + LR	W2V-Avg + Dense	LSTM (pre)	LSTM (scratch)
0	can you remind me to not call my mom	N/A	calendar_set	general_joke	general_dontcare	email_sendemail
1	is it going to be sunny or rainy tomorrow	N/A	weather_query	calendar_query	lists_createoradd	social_post
2	find a flight from new york to london but not ...	N/A	general_negate	transport_query	alarm_set	cooking_recipe

Nhận xét:

- Các câu có phủ định ('not') hoặc mệnh đề phụ ('or', 'but not through ...') thường khiến mô hình TF-IDF và W2V trung bình dự đoán sai.
- LSTM (đặc biệt là pre-trained) có xu hướng hiểu ngữ cảnh tốt hơn, vì trạng thái ẩn theo chuỗi giúp mô hình nhận diện được mối quan hệ giữa các từ.
- Nếu LSTM (pre-trained) chính xác hơn bản học từ đầu, điều đó chứng tỏ embedding Word2Vec cung cấp ngữ nghĩa ban đầu hữu ích, g